

EFFICIENT ANY-TARGET BACKDOOR ATTACK WITH PSEUDO POISONED SAMPLES

Bin Huang, Zhi Wang

Tsinghua Shenzhen International Graduate School, Tsinghua University, Shenzhen, China

ABSTRACT

Deep neural networks present their potential vulnerabilities under backdoor attacks. They have a satisfactory performance for benign users with clean samples but will get malicious outputs when inputs are attached with the backdoor trigger. Current backdoor attacks usually target only one single class, making them not robust to defenses against this characteristic. In this work, we propose an invisible and input-dependent any-target attack that targets all the labels simultaneously. Specifically, we train the model to be attacked together with the image steganography model by encoding the one-hot encodings into the images. The novel pseudo poisoned samples are introduced to improve the effectiveness of our attack. Experimental results show that our method is both effective and efficient on several datasets and is robust to existing defenses. Our code is available at <https://github.com/huangbinary/any-target-attack-with-PPS>.

Index Terms— Backdoor Attack, Any-Target Attack, Steganography, Pseudo Poisoned Samples

1. INTRODUCTION

Deep neural networks have been applied in a variety of applications due to their powerful ability of data representation. However, the rigorous necessity of massive training data, expensive computing resources and long training time discourage the entities that lack these conditions. As an alternative, they may turn to third-parties for training their models or using the pre-trained ones directly, which brings new surface of potential security threats. The backdoor attack is a well-studied one of them.

BadNets [1] is the first proposed backdoor attack on neural networks, where the attacker poisons a portion of training samples by attaching a pre-defined pattern (termed *trigger*) on them and changes the corresponding labels to a target one. Models trained on this partially modified training set will be embedded with hidden backdoor, predicting the target label when the trigger shows up while behaving normally otherwise. Studies on backdoor attack usually focus on its effectiveness or stealthiness. To improve the effectiveness, the attacker tries to find out better triggers to increase the attack success rate [2, 3]. While for the stealthiness, as existing attacks often rely on triggers that are visi-

ble and static, making them easily noticeable visually or detectable algorithmically, many invisible [3, 4], dynamic [5, 6] or input-dependent [5, 7, 8] triggers have been proposed continually. What’s more, backdoor attack without label modification, which is called clean-label attack [9, 10, 11], can also enhance the stealthiness.

From the perspective of the target labels, we argue that targeting multiple labels simultaneously (multi-target attack) will make the attack harder to detect. There is a kind of multi-target attack named all-target attack in current studies [1, 5], where the attacker aims to misclassify poisoned samples as their next labels, i.e. classify samples in class c as class $c + 1$.

In this paper, we investigate a more aggressive attack called any-target attack, where the attacker generates poisoned samples of any target classes and leads the attacked model to predict them as the specified classes. Some existing single-target attacks [1, 12, 13] can be easily extended to the any-target attacks by mapping distinct triggers to different labels. But the process is not straightforward and the extended attacks are inefficient and ineffective experimentally (in Sec. 3.2). A few attacks [6, 14] are applicable to the any-target attack paradigm, but they are visible and input-independent.

To overcome above shortcomings, inspired by prior works on image steganography [15, 16] and the steganography based backdoor attacks [7, 17], we propose an invisible and input-dependent any-target attack by training the model to be attacked together with the image steganography model. The target label is specified by its corresponding one-hot encoding which is then encoded into the benign input image to produce a poisoned one. As thus, the tedious process of selecting distinct triggers in any-target attack can be efficiently implemented by just offering the one-hot encodings. What’s more, the poisoned images are invisible and input-dependent owing to the encoder of the image steganography model. In the end, to encourage the attacked model to better learn the mapping from the triggers to the target labels, we introduce the novel pseudo poisoned samples which are also poisoned samples generated by the encoder but with their labels unchanged, ensuring the effectiveness of the any-target attack.

We evaluate our attack on four datasets MNIST, CIFAR-10, GTSRB and CelebA. Experiments show that on every dataset we achieve high attack success rate of over 99% and similar benign accuracy with the benign model, even under a rather low attack ratio. The triggers added to the poisoned

images are visually imperceptible. The proposed attack can also bypass several typical defense methods. Ablation studies reveal the importance of the novel pseudo poisoned samples to the success of our attack whose potential of attacking larger number of classes has been verified as well.

Our main contributions are listed as follows:

- (1) We explore an invisible and input-dependent any-target backdoor attack to make it more stealthy to human inspection or algorithmic detection.
- (2) We introduce novel pseudo poisoned samples to improve the performance of our attack.
- (3) Extensive experiments demonstrate the effectiveness and efficiency of the proposed method and its robustness to several defense methods.

2. THE PROPOSED ATTACK

2.1. Problem Definition

Definition of Traditional Backdoor Attack. Given a clean training set $\mathcal{D} = \{(x, y)\}$, where x is an input image and y is the corresponding label, an attacker applies a backdoor injection function \mathcal{B} and a label mapping function \mathcal{M} on part of \mathcal{D} to generate a poisoned training set $\mathcal{T} = \{(x_t, y_t)\} = \{(\mathcal{B}(x), \mathcal{M}(y)) | (x, y) \in \mathcal{D}', \mathcal{D}' \subset \mathcal{D}\}$, where x_t is the poisoned sample and y_t is its poisoned label. For example, BadNets [1] defines \mathcal{B} as adding a static small patch on the fixed position of an image as follows:

$$\mathcal{B}(x) = x \odot (1 - m) + p \odot m, \quad (1)$$

where p is a pattern, m a blending mask and \odot pixel-wise multiplication. In prior works, \mathcal{M} is often a constant function $\mathcal{M}(y) = t$ (single-target attack), where t is the target label, or a shifting function $\mathcal{M}(y) = (y + 1) \bmod Y$ (all-target attack), where Y is the class number. The attack ratio is defined as the ratio of the poisoned sample volume to the training sample volume, i.e. $\alpha = |\mathcal{T}|/|\mathcal{D}|$. After poisoning the training set, the attacker injects a backdoor into the classification model by training it over \mathcal{T} and the non-poisoned part of \mathcal{D} , i.e. $\mathcal{T} \cup \mathcal{D} \setminus \mathcal{D}'$. The goal of the attacker is to make targeted classification if the backdoor trigger is present at test-time samples while without harming its classification accuracy on benign test samples.

Definition of Any-Target Backdoor Attack. We refer to our attack paradigm as any-target attack to distinguish it from the all-target attack. Different from the definition of previous attack where the backdoor injection function \mathcal{B} is only conditioned on the input image, in any-target attack, \mathcal{B} takes as input an image as well as a target class t , then it generates a poisoned sample via $x_t = \mathcal{B}(x, t)$. While the mapping function \mathcal{M} just outputs t used in the previous generating process. Note that t can be any class instead of a fixed one. By encoding the class information into the image, there will be some perturbation in part of the pixel values. We can treat it as the trigger added to the sample for backdoor effect. It should

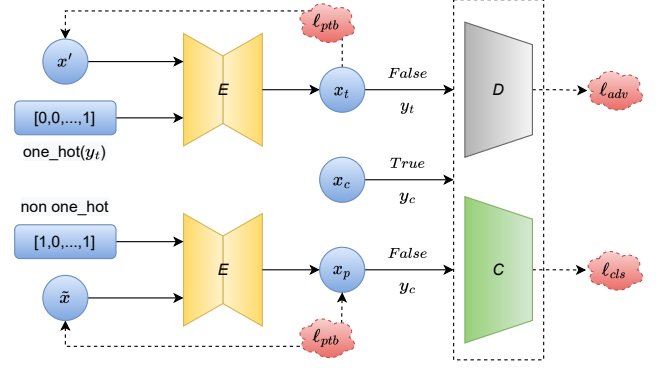


Fig. 1. The overall workflow of our proposed method. Three modes of the training process are illustrated. E , C and D are the encoder, classifier and discriminator respectively. The labels after x are ground truth for D and C respectively.

be noted that the perturbed pixel number or intensity must be small enough for invisibility consideration.

Attacker's Capacities. We consider one commonly used threat model by existing backdoor attack studies. The attacker has full knowledge of the training dataset and training algorithm. He can modify the training data as he want and intervene the training process for desired purpose. The infected model is then distributed for users to download.

Attacker's Goals. After injecting the backdoor, it is expected that two main goals are achieved. To be specific, on clean images, the attacked model shows similar classification performance with that not attacked, while on the generated samples by \mathcal{B} , the model predicts the target class t which is encoded in the input with high probability.

2.2. Proposed Method

Overview. The overall framework of our proposed attack is illustrated in Fig. 1. We take the image steganography model Hidden [15] as our base. Specifically, Hidden comprises four main components: an encoder, a noise layer, a decoder and a discriminator. We keep its encoder and discriminator as our *encoder* E and *discriminator* D , while replace its decoder with the *classifier* C to be attacked, as our purpose is to classify the encoded image as the target class rather than recover the encoded message from it.

Given a dataset \mathcal{D} with N samples, we randomly take α of them as the poisoned set \mathcal{D}' . The encoder E receives an image x' from \mathcal{D}' and a **random** target label y_t and then produces a poisoned image x_t which is encoded with y_t . Note that any vector can be used to represent the target label if only they are one-to-one mapped. In our method, we choose the one-hot encoding for simplicity. The classifier C predicts an image x_c from \mathcal{D} as its ground truth class y_c and the poisoned image x_t as its encoded class y_t . At the same time, the discriminator D tries to predicts the probability that the input image, either x_c or x_t , is not poisoned.

The poisoned image should have no obvious visual difference with the corresponding benign one. We limit the variation of the label encoding process with a *perturbation loss* ℓ_{ptb} , the ℓ_2 distance between x' and x_t , and a *adversarial loss* ℓ_{adv} , the binary cross entropy (BCE) that aims to fool the discriminator. For the attacked classifier, we impose a *classification loss* ℓ_{cls} using the cross entropy (CE) to ensure the expected prediction. The discriminator incurs *adversarial loss* from both x' and x_t , which reflects the ability to distinguish the poisoned images.

Pseudo Poisoned Samples. Owing to the limited number of poisoned samples compared with that of clean samples, coupled with the large number of targeted classes, the attacked model can hardly well learn the representation of clean samples and poisoned samples simultaneously. To this end, we introduce the *pseudo poisoned samples* to enhance the representation learning while keeping the accuracy on benign samples, making the attack more effective and more efficient.

A pseudo poisoned samples is generated by the encoder via encoding into it some vector dubbed *pseudo encoding*. Since the one-hot encodings are taken for representing the target labels, we randomly choose one of the rest non one-hot encodings as the input vector of E , that is, there are zero or more than one 1 in it. It should be noted that the length of the pseudo encoding equals to a one-hot encoding's. The difference between poisoned samples and pseudo poisoned samples lies in their ground truth labels.

We first randomly take the same amount, i.e. ratio α , of samples from \mathcal{D} as the pseudo poisoned set $\tilde{\mathcal{D}}$. Then, an image \tilde{x} from $\tilde{\mathcal{D}}$ is sent to E in company with a pseudo encoding, resulting a pseudo poisoned sample x_p . When calculating the classification loss, the ground truth label remains unchanged. Other parts of the workflow of \tilde{x} are identical with that of x' . After introducing the pseudo poisoned samples, we improve the ability of the attacked model to distinguish poisoned samples from others. As is shown in Sec. 3.2, the attack performance improves significantly.

Training Process. The three modes of the training process are summarized as: (1) **Clean mode.** The classifier C has to correctly recognize x_c . (2) **Attack mode.** When x_t is encoded with a target label, the attack is activated, making C output wrong prediction y_t . (3) **Pseudo attack mode.** Although encoded with a similar vector, x_p can still not fool C .

In one epoch, we first jointly train the encoder E and classifier C by minimizing the loss \mathcal{L}_1 in Eq. (2) through stochastic gradient descent. And then the discriminator D is trained likewise by minimizing the loss \mathcal{L}_2 in Eq. (3) as follows:

$$\begin{aligned}\mathcal{L}_1 &= \mathbb{E}_{(x,y) \in \mathcal{D}} [\lambda_{cls} \ell_{cls}(C(x), y)] \\ &\quad + \mathbb{E}_{(x,y) \in \mathcal{D}' \cup \tilde{\mathcal{D}}} [\lambda_{cls} \ell_{cls}(C(E(x)), \mathcal{M}(y)) \\ &\quad + \lambda_{ptb} \ell_{ptb}(x, E(x)) + \lambda_{adv} \ell_{adv}(D(E(x)), 1)], \quad (2) \\ \mathcal{L}_2 &= \mathbb{E}_{x \in \mathcal{D}' \cup \tilde{\mathcal{D}}} [\ell_{adv}(D(x), 1) + \ell_{adv}(D(E(x)), 0)], \quad (3)\end{aligned}$$

where λ_{cls} , λ_{ptb} and λ_{adv} are hyperparameters that balance

each loss. Note that the two items in \mathcal{L}_2 share equal weight.

3. EXPERIMENTAL RESULTS

3.1. Experimental Settings

Datasets and Models. We conduct our attack experiments on four datasets: MNIST, CIFAR-10, GTSRB and CelebA. The classifier is a 4-layer convolutional neural network for MNIST following previous studies [5] and Pre-activation Resnet-18 for others. The encoder and discriminator are kept consistent with Hidden [15] except for the input layer. To verify the effectiveness on larger number of classes, we randomly select a subset of 200 classes from ImageNet.

Training Setup. We adopt the SGD optimizer with the initial learning rate set as 0.01, which drops to 0.1 times after every 100 epochs. The batch size and maximum epoch are 128 and 300 respectively. We set the attack ratio α as 0.1 and decrease it by 0.02 each time for comparison until 0.02. λ_{cls} , λ_{ptb} and λ_{adv} are set to 1, 0.7 and 0.001. Experiments are conducted on NVIDIA GeForce RTX 3090 GPUs.

Baseline Selection. We extend BadNets [1], Blended [18] and ISSBA [7] to any-target attack. To be specific, N distinct triggers are elaborately selected for all N classes. We compare our method with those attacks as well as the benign model without backdoor injection.

Evaluation Metrics. We use benign accuracy (ACC) and attack success rate (ASR) for evaluation. They are defined as the accuracy on benign and poisoned samples respectively.

3.2. Main Results


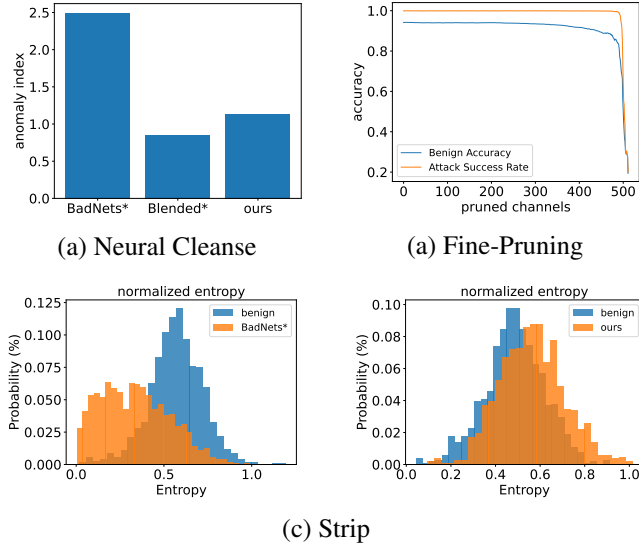
Performance under Different Attack Ratios. Tab. 1 reports the comparison of different attacks. As can be seen, under $\alpha = 0.1$, our attack can achieve a high ASR approaching to 100% on all datasets, better than the others, while keeping ACC very close to the benign model. On the contrary, extended BadNets or Blended can only have comparable high ASR on part of the datasets. We also experiment on the extended ISSBA, but result a failure attack where ASR equals to the probability of random guess, i.e. the reciprocal of the class number. Besides, even under more lower α , the ASR of our attack still stays at a high level of over 99%, while that of other attacks drops significantly, revealing the effectiveness and efficiency of ours. Note that there is slight reduction on the ACC of our method when α decreases to 0.2, which means it focuses more on learning the poisoned samples.

Visualization of Poisoned Images. Fig. 2 presents an example of the poisoned images generated by our method. The poisoned image brings satisfactory visual invisibility, making it look natural to human inspection and reflecting the stealthiness of our attack.

Resistant to Backdoor Defenses. We test our approach against three defenses: Neural Cleanse [19] (NC), Fine-Pruning [20] and Strip [21]. NC reverse engineers candidate

Table 1. Attack performance (%) of different methods. The best result is in bold. * means the corresponding extended attack.

| attack ↓ | $\alpha \rightarrow$ | MNIST | | | CIFAR10 | | | GTSRB | | | CelebA | | |
|----------|----------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| | | 0.1 | 0.06 | 0.02 | 0.1 | 0.06 | 0.02 | 0.1 | 0.06 | 0.02 | 0.1 | 0.06 | 0.02 |
| benign | ACC | 99.58 | | | 94.74 | | | 99.20 | | | 80.26 | | |
| BadNets* | ACC | 99.47 | 99.40 | 99.44 | 93.52 | 93.82 | 94.04 | 98.64 | 99.03 | 98.60 | 79.94 | 79.29 | 79.09 |
| | ASR | 95.52 | 95.18 | 89.78 | 48.84 | 49.64 | 48.35 | 94.94 | 93.03 | 59.77 | 77.41 | 78.55 | 77.80 |
| Blended* | ACC | 99.59 | 99.60 | 99.50 | 93.68 | 94.00 | 94.21 | 98.62 | 98.73 | 98.95 | 79.91 | 79.57 | 79.27 |
| | ASR | 99.97 | 99.92 | 97.98 | 93.86 | 88.25 | 59.57 | 90.58 | 82.17 | 56.00 | 99.66 | 99.11 | 97.12 |
| ours | ACC | 99.60 | 99.65 | 99.47 | 94.18 | 94.02 | 92.21 | 98.54 | 98.91 | 99.04 | 79.44 | 79.31 | 78.76 |
| | ASR | 100.0 | 100.0 | 100.0 | 99.97 | 99.84 | 99.45 | 99.26 | 99.01 | 99.81 | 99.70 | 99.89 | 99.58 |

**Fig. 2.** An example of the visualization of poisoned images.**Fig. 3.** Experiments of three defenses against our attack.

triggers for each class and finds out significantly small one as the target. It calculates an *anomaly index* which indicates a backdoor when greater than 2. As shown in Fig. 3(a), our attack successfully bypass NC. When pruning the channels on the penultimate layer increasingly, our attack retains its high ASR even when almost all the channels are pruned as in Fig. 3(b), showing the resistance to pruning-based defense. Strip determines candidate poisoned samples by calculating the entropy of the average prediction on the samples generated by imposing various pre-prepared clean samples on the

Table 2. Performance (%) comparison of our attack with and without pseudo poisoned samples (PPS).

| attack ↓ | CIFAR10 | | GTSRB | | CelebA | |
|----------|--------------|--------------|--------------|--------------|--------------|--------------|
| | ACC | ASR | ACC | ASR | ACC | ASR |
| w/ both | 94.18 | 99.97 | 98.54 | 99.26 | 79.44 | 99.70 |
| w/o PPS | 88.38 | 76.68 | 18.27 | 97.85 | 75.17 | 97.88 |

testing samples. Smaller entropy means easier detection of the poisoned samples. But in our attack, as shown in Fig. 3(c), the entropy is close to, even slightly larger than that in the benign model, which represents robustness to Strip. All the results on Fig. 3 are on CIFAR-10 and results on other datasets are similar.

Necessity of Pseudo Poisoned Samples. We carry out a variation of our attack, which consists of only the first two modes in the training process, that is, without pseudo attack mode. As we can see in Tab. 2, after removing the pseudo poisoned samples, both ACC and ASR drop dramatically, which demonstrates the necessity of them.

Potential of Target Class Numbers. We perform several attacks on a subset of 200 classes from ImageNet for every additional 20 classes. Experimental results indicate that the ASR drops a little as the class number increases when it is larger than 140, but still reaches about 90% when it is 200.

4. CONCLUSION

In this work, we propose an invisible and input-dependent any-target attack from the perspective of the target labels to improve the stealthiness of the attack. By training the model to be attacked together with the encoder of an image steganography model, the specified target label can be successfully encoded into the image as invisible trigger. With the assistance of the novel pseudo poisoned samples, the attacked model can well learn the representation of different samples. Experimental results show the effectiveness and efficiency of our method even under a very low attack ratio.

5. REFERENCES

- [1] Tianyu Gu, Brendan Dolan-Gavitt, and Siddharth Garg, “Badnets: Identifying vulnerabilities in the machine learning model supply chain,” *arXiv preprint arXiv:1708.06733*, 2017.
- [2] Yingqi Liu, Shiqing Ma, Yousra Aafer, Wen-Chuan Lee, Juan Zhai, Weihang Wang, and Xiangyu Zhang, “Trojaning attack on neural networks,” 2017.
- [3] Khoa Doan, Yingjie Lao, Weijie Zhao, and Ping Li, “Lira: Learnable, imperceptible and robust backdoor attacks,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021, pp. 11966–11976.
- [4] Erwin Quiring and Konrad Rieck, “Backdooring and poisoning neural networks with image-scaling attacks,” in *2020 IEEE Security and Privacy Workshops (SPW)*. IEEE, 2020, pp. 41–47.
- [5] Tuan Anh Nguyen and Anh Tran, “Input-aware dynamic backdoor attack,” *Advances in Neural Information Processing Systems*, vol. 33, pp. 3454–3464, 2020.
- [6] Ahmed Salem, Rui Wen, Michael Backes, Shiqing Ma, and Yang Zhang, “Dynamic backdoor attacks against machine learning models,” in *2022 IEEE 7th European Symposium on Security and Privacy (EuroS&P)*. IEEE, 2022, pp. 703–718.
- [7] Yuezun Li, Yiming Li, Baoyuan Wu, Longkang Li, Ran He, and Siwei Lyu, “Invisible backdoor attack with sample-specific triggers,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021, pp. 16463–16472.
- [8] Le Feng, Sheng Li, Zhenxing Qian, and Xinpeng Zhang, “Stealthy backdoor attack with adversarial training,” in *ICASSP 2022-2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2022, pp. 2969–2973.
- [9] Ali Shafahi, W Ronny Huang, Mahyar Najibi, Octavian Suci, Christoph Studer, Tudor Dumitras, and Tom Goldstein, “Poison frogs! targeted clean-label poisoning attacks on neural networks,” *Advances in neural information processing systems*, vol. 31, 2018.
- [10] Aniruddha Saha, Akshayvarun Subramanya, and Hamed Pirsiavash, “Hidden trigger backdoor attacks,” in *Proceedings of the AAAI conference on artificial intelligence*, 2020, vol. 34, pp. 11957–11965.
- [11] Alexander Turner, Dimitris Tsipras, and Aleksander Madry, “Label-consistent backdoor attacks,” *arXiv preprint arXiv:1912.02771*, 2019.
- [12] Xueluan Gong, Yanjiao Chen, Qian Wang, Huayang Huang, Lingshuo Meng, Chao Shen, and Qian Zhang, “Defense-resistant backdoor attacks against deep neural networks in outsourced cloud environment,” *IEEE Journal on Selected Areas in Communications*, vol. 39, no. 8, pp. 2617–2631, 2021.
- [13] Nan Zhong, Zhenxing Qian, and Xinpeng Zhang, “Imperceptible backdoor attack: From input space to feature representation,” *arXiv preprint arXiv:2205.03190*, 2022.
- [14] Xinzhe Zhou, Wenhao Jiang, Sheng Qi, and Yadong Mu, “Multi-target invisibly trojaned networks for visual recognition and detection,” in *IJCAI*, 2021, pp. 3462–3469.
- [15] Jiren Zhu, Russell Kaplan, Justin Johnson, and Li Fei-Fei, “Hidden: Hiding data with deep networks,” in *Proceedings of the European conference on computer vision (ECCV)*, 2018, pp. 657–672.
- [16] Matthew Tancik, Ben Mildenhall, and Ren Ng, “Stegastamp: Invisible hyperlinks in physical photographs,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 2117–2126.
- [17] Shaofeng Li, Minhui Xue, Benjamin Zi Hao Zhao, Haojin Zhu, and Xinpeng Zhang, “Invisible backdoor attacks on deep neural networks via steganography and regularization,” *IEEE Transactions on Dependable and Secure Computing*, vol. 18, no. 5, pp. 2088–2105, 2020.
- [18] Xinyun Chen, Chang Liu, Bo Li, Kimberly Lu, and Dawn Song, “Targeted backdoor attacks on deep learning systems using data poisoning,” *arXiv preprint arXiv:1712.05526*, 2017.
- [19] Bolun Wang, Yuanshun Yao, Shawn Shan, Huiying Li, Bimal Viswanath, Haitao Zheng, and Ben Y Zhao, “Neural cleanse: Identifying and mitigating backdoor attacks in neural networks,” in *2019 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2019, pp. 707–723.
- [20] Kang Liu, Brendan Dolan-Gavitt, and Siddharth Garg, “Fine-pruning: Defending against backdooring attacks on deep neural networks,” in *International Symposium on Research in Attacks, Intrusions, and Defenses*. Springer, 2018, pp. 273–294.
- [21] Yansong Gao, Change Xu, Derui Wang, Shiping Chen, Damith C Ranasinghe, and Surya Nepal, “Strip: A defence against trojan attacks on deep neural networks,” in *Proceedings of the 35th Annual Computer Security Applications Conference*, 2019, pp. 113–125.